

Lecture 2 - Data Models

What's happening today

- Main paper: "What goes around comes around" (2005), Stonebraker (mainly), according to Joe.
- Second paper: Revisiting that via "...and around" (2025), Stonebraker and Pavlo
- Both papers are very opinionated (you can think of them as "thought pieces")
 - **do not** believe everything as written!
 - heavy bias towards systems the authors have written!
- Third (optional) paper: "A relational model of data for large shared data banks"
 - yes, this is the Codd paper
- Goal - revisit explorations of various data models over time
 - 1960s–2005; 2005–2025
 - what were the lessons from these eras - on the data model side
- Let's start by talking about the first paper, which was part of the Redbook
 - Main motivation of that paper was that all data models were revisiting arguments in the hierarchical/network data model and the relational data model
- First, some concepts from the Codd paper:
 - Physical data independence & logical data independence
 - Physical: applications continue to work as physical organization changes
 - Logical: applications continue to work as logical organization (schema) changes

Various data model eras

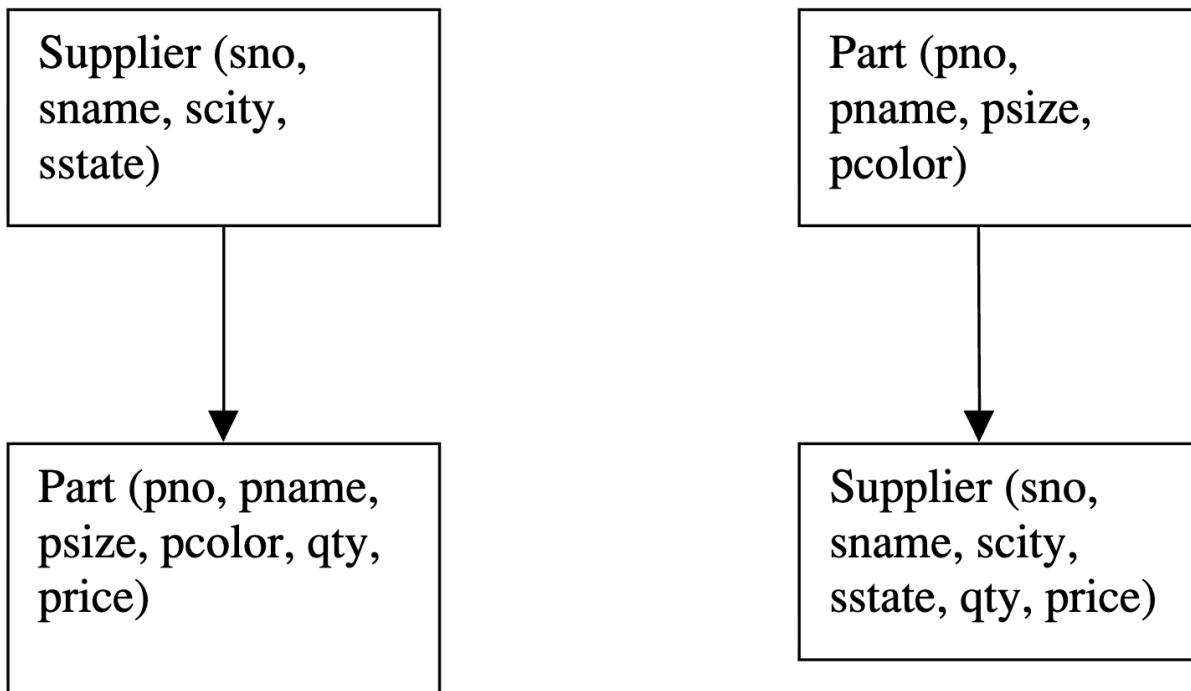
- Hierarchical (IMS): late 1960's and 1970's
- Network (CODASYL): 1970's
- Relational: 1970's and early 1980's
- Entity-Relationship: 1970's
- Extended Relational: 1980's
- Semantic: late 1970's and 1980's
- Object-oriented: late 1980's and early 1990's
- Object-relational: late 1980's and early 1990's

Example setup

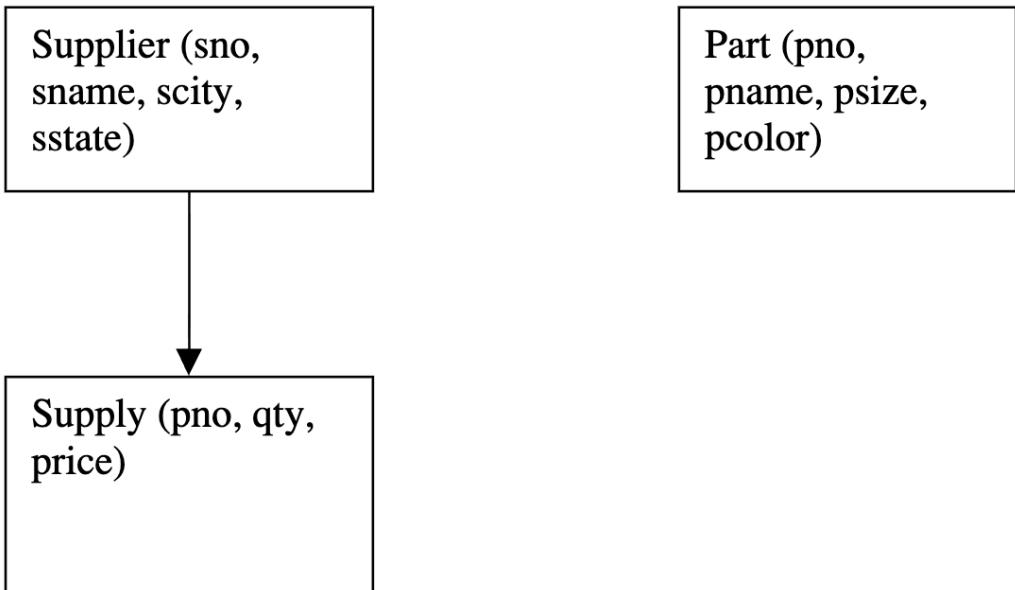
- Supplier (sno, sname, scity, sstate)
- Part (pno, pname, psize, pcolor)
- Supply (sno, pno, qty, price)

IMS (hierarchical) era (1960s)

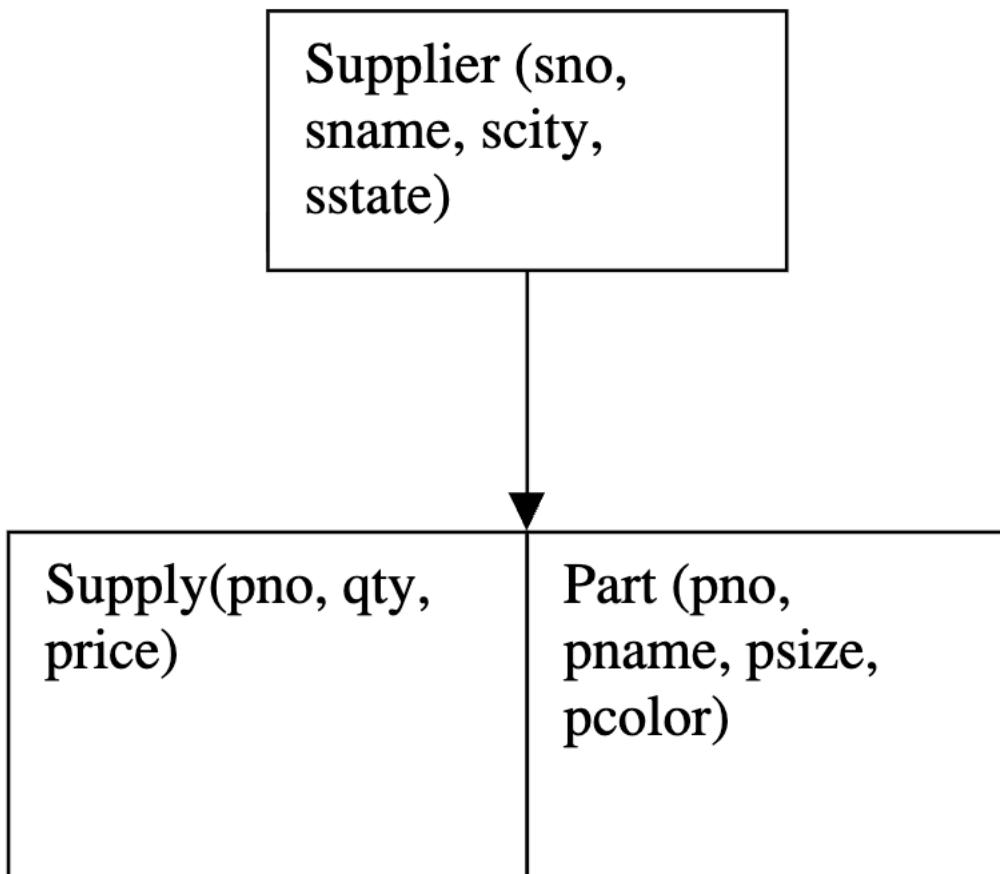
- Data Model:
 - Records have fields and a key (a set of fields)
 - Record types (schemas) are hierarchical:
 - arranged in a rooted ordered (will see why this matters) tree, where each record type has a parent record type
- One approach - where Supply gets "squished" into child node or alternative



- Main issues with this approach:
 - Redundancy:
 - On LHS, part info is repeated for each supplier who supplies the part
 - Redundancy is not great because it raises the danger of inconsistencies
 - Existence: A part can't exist without a supplier (same for other)
 - Some queries are easier than others - depending on nesting order
- Query Execution:
 - Record-at-a-time
 - Retrieval of records done by using hierarchical sequence key: a concatenation of keys along the path to the root
 - Users need to figure out where to start and how to traverse
 - Expectation that traversal happens depth-first, left-to-right
 - Essentially following physical pointers, or sequential scan
- On Data Independence:
 - Violates physical data independence because the program could expect that the data is laid out in a certain way, e.g., pointers from parents to children.
 - Some amount of logical data independence because you can extend tree by adding additional roots
- To avoid redundancy, one can do the following:



- But the challenge is that combining information across "trees" is impossible
 - so mechanism to do so is to physically stitch things together through pointers,
 - from the programmer's standpoint this looks like this, even though there's only one copy of each part, which is reachable from each supply

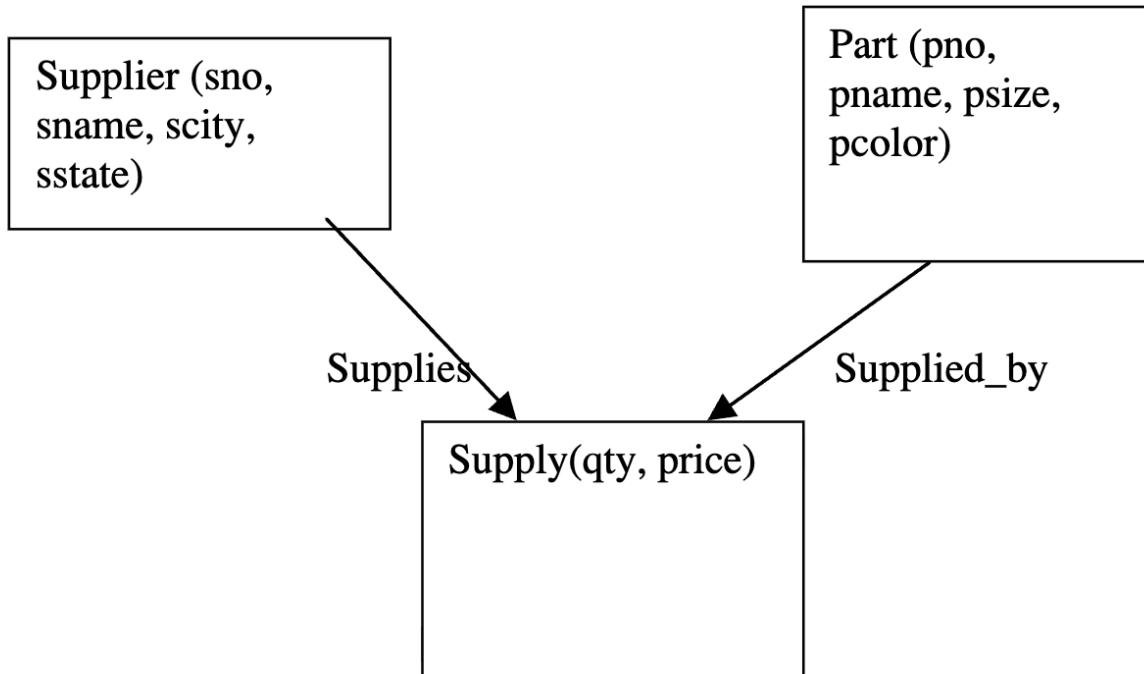


- Overall, tree structured data models have various downsides

- don't provide any real physical or logical data independence.
 - can't really reorganize the data
- record-at-a-time imperative query specification approach requires users to make tradeoffs:
 - should they start at the root and traverse down following pointers, or:
 - jump to a node and then walk down, etc.
 - no real query optimization across these alternatives - all manual
- the representations are restrictive - combining information across entities isn't possible unless they are hierarchically organized within each other
- schema defines both logical relationships and physical navigation paths
 - schema equals access methods

CODASYL (network) era (1970s)

- This got Charles Bachmann the Turing award
- Main issue with IMS is that each record type has a single parent
 - Record types as in IMS, but now a network rather than a tree



- Data Model
 - Basically, data model: record types + sets
 - Supplies and Supplied_by edges (these are called "sets") indicate that for each parent, there is a relationship with 0 or more children.
 - Each set, however, can only support a two-way relationship, not multi-way
 - Eg sale - salesperson, customer, item
- Query Execution:
 - Still record-at-a-time` `Find Supplier (SNO = 16) Until no-more { Find next Supply record in Supplies Find owner Part record in Supplied_by Get current record -check for red- }
 - Essentially following pointers between parents and children
 - But now, even more complicated, since one needs to keep track of what all one has visited in the context of each individual loop and also globally
 - In IMS, the programmer needs to only keep track of where they are currently, and their parent
 - Here, one may need to go down, sideways, up, and down again ...

- Essentially, a user is implementing a form of complicated directed nested loop, manually - keeping track of where they are for each step
 - Bachmann: "navigating in hyperspace"
- Independence:
 - Because of these pointers, no physical data independence
 - Similarly, programs require the network schema, so no logical data independence

Relational Era (1970s and 1980s)

- From paper "*the driver ... was the fact that IMS programmers were spending large amounts of time doing maintenance on IMS applications, when logical or physical changes occurred.*"
- Three maxims (from paper)
 - *Store the data in a simple, flat data structure (tables)*
 - flexible also to represent multi-way relationships
 - *Access it through a high level set-at-a-time DML*
 - no imperative nested loops to traverse data
 - *No need for a physical storage proposal*
 - database system decides how to store
 - this means no pointers, no expectation that records can be located by following certain paths
- Truly declarative - users are not manually optimizing their programs, systems can figure it out
- Led to "the great debate" between Codd and Bachmann at SIGMOD'74
 - Codd: CODASYL is too complex, too little physical independence, nested-loop based programming is hard to optimize
 - Bachmann: Can never be efficient
- By '76, INGRES and System R had shown that efficient implementations were possible; relational model came with physical data independence anyway.
 - INGRES also introduced the notion of views, adding to logical data independence
- By '84 IBM prioritized relational systems - ending the "great debate"

ER Model (1970s)

- Originally proposed by Peter Chen as an alternative to relational, CODASYL, hierarchical data models - and trying to unify them
- Data model
 - entities + relationships, each with attributes



- One way it was "better" than relational is that it made relationships a first-class citizen
- Did not actually have a proper query execution proposal
- Eventual use-case
 - DB design (translate into relational), rather as a serious data model proposal - no serious system supported it
 - Alongside, DB theory folks developed theories of normalization (including Codd, who introduced normalization in the first place)
 - 2NF, 3NF, BCNF, 4NF

- Stonebraker disses on normalization "*Functional dependencies are too difficult for mere mortals to understand.*"
- But still very useful to understand where redundancy can occur in a schema
 - Even if we choose to not do anything about it

Relational Extensions (early to mid 1980s), Semantic data models (1980s)

- Relational ++
 - Supporting set oriented attributes:
 - e.g., Part(pno, pname, available_colors = {red, blue, green})
 - Avoids joins but breaks atomic attribute property
 - Foreign keys via tuple pointers
 - e.g., Supply(PT: Part, SR: Supplier, qty, price), PT, SR - pointers to tuples in those tables
 - Like CODASYL's pointers
 - Sub-classing (specialization/generalization)
 - Didn't really catch on until OO languages took off
- Semantic data models
 - Added in nested records
 - Multiple inheritance
- Why didn't this take off?
 - Mainly this wasn't where the pain was: RM could handle it
 - Not enough added functionality, most of the pain was in performance

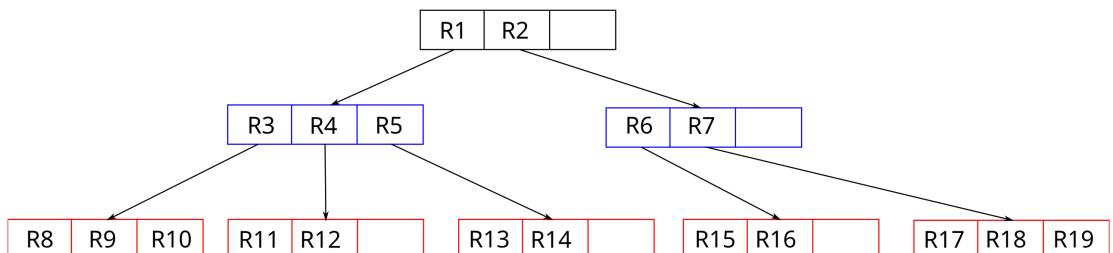
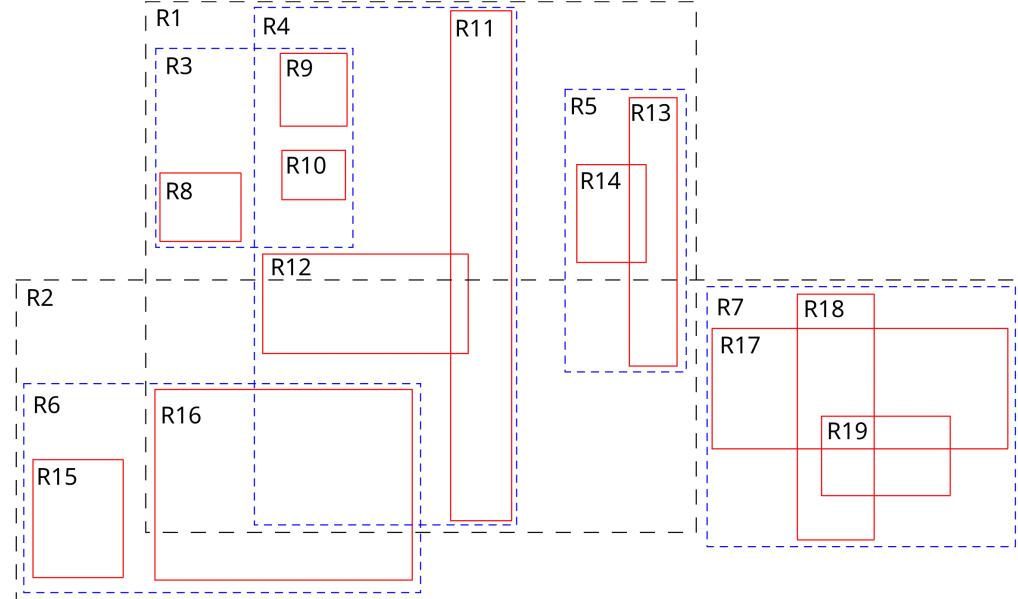
OODB (1980s to 1990s)

- C++ on the rise, with object orientation
- Main problem they wanted to solve:
 - impedance mismatch between programming languages and databases
 - also wanted support for complex objects (e.g., CAD)
- Made PL objects = database objects
 - Declare a variable -> gets persisted
 - Persistent Part p;
 - p.color = "red"; // persisted
 - Early versions of ORMs
- Why didn't this take off?
 - People wanted to not be tied to a programming language
 - Needed buy in from PL folks
 - Ends up being network/hierarchical data model again

OR (1980s and 1990s)

- Major motivation was GIS (map data)
 - Want to be able to support range queries (lat long), without having separate indexes for lat and long - as this is pretty inefficient
 - Need for complex data types and support for indexes on complex data types
 - Want all of this as part of SQL
- Core idea: make DB extensible
 - User defined types, operators, functions, access methods
 - e.g., intersection operator "WHERE point spatially_in "x0, x1, y0, y1""
 - internally calls UDF - user defined function

- Basically database engine needs more flexibility to support these complex data types
- Outcomes:
 - Postgres was the major successful system of this era
 - New indexes become commonplace
 - R-tree, Quad-Tree
 - Detour to present R-trees
 - Height balanced tree (like B-tree)
 - Each node has between m to M pointers
 - With each pointer, we store MBR: minimum bounding rectangle
 - For search, we will go down each MBR that overlaps with rectangle (maybe multiple search paths in worst case)
 - For insertion, go down MBR that will expand the least at each step
 - Find leaf where the object fits, and if it does, great
 - If not, split, and then propagate split upwards (just like B+tree)



- Another innovation: stored procedures
 - Define multi-SQL queries that can be treated as a unit - essentially a predefined UDF
- However, the extensibility, while really helpful for having relational databases adapt to new requirements, wasn't the major win.
 - paper: "*the major win in Postgres was UDTs and UDFs; the OO constructs were fairly easy and fairly efficient to simulate on conventional relational systems*"
 - UDFs bring code to data, as opposed to data to code - reducing data transfer
 - specialized programs that can't be expressed in SQL
 - To this day, UDFs are essential to support important "external" primitives; think any libraries, ML, LLMs, specialized functions for each vertical...

- Another way to think of this is that it was an attempt towards adding imperative PL constructs into a declarative system
- from paper: "*The major benefits of OR is two-fold: putting code in the data base (and thereby blurring the distinction between code and data) and user-defined access methods.*"

Semi-structured data (2000s)

- Two main ideas
 - Schema-less, instead self-describing format
 - Nested, hierarchical model
- Canonical setup
 - XML (now JSON)
 - `<person><name>Mike Stonebraker</name><age>85</age></person>`
 - Set of attributes can be different; could be nested
 - Paper walks through various use cases, arriving at one compelling one: resumes
 - Moving it closer to relational data - XML Schema, DTDs
 - Paper's prediction "*XML Schema will fail because of excessive complexity*"
- Parallels to IMS and CODASYL
 - Hierarchical like IMS, with same issues of redundancy etc.
 - Some CODASYL features too
 - Set based attributes, like SDM
 - From paper: "*The current proposal is now a superset of the union of all previous proposals. I.e. we have navigated a full circle.*"
- Some deviation from IMS/CODASYL
 - instead of users looping through, and using pointers
 - Instead use set-at-a-time query language, Xquery (or XPath)
- Implementation-wise: Challenges in supporting even a simple variant: moving from atomic type to union type, is difficult
 - what would index look like - require hybrid index
 - operators require custom implementations - need to account for each type interacting with each type of the other "side"
- What has stood the test of time:
 - semi-structured data is an excellent interchange format

Main takeaways from paper

- Only two good ideas from the entire set of data model proposals:
 - UDFs
 - Schema last - but only for data interchange

(Didn't actually cover the following in class)

... and around

- Paper from 2024 - recaps a number of advances from the 2000s and 2010s
 - Many of these are not really data models but different system architectures
- Their conclusion is similar "*that the relational model with an extendable type system (i.e., object-relational) has dominated all comers, and nothing else has succeeded in the marketplace.*""
 - Generally true but there are specialized systems that work really well in certain settings
 - What is also true that the NoSQL movement of the late 2000s has come full circle to embrace RM and SQL

- Mostly rejected RM as being not good enough performance-wise

Various new proposals

1. Map Reduce Systems
2. Key-value Stores
3. Document Databases
4. Wide-Column
5. Text Search Engines
6. Array Databases
7. Vector Databases, and
8. Graph Databases.

1. Map Reduce

- Parallel programming framework, based on low-level primitives map and reduce
 - One way they describe it is as: `SELECT map() FROM crawl_table GROUP BY reduce()` - but `map()` is done first
- Not really a data model proposal, more a programming abstraction proposal but suffers from similar issues
 - Moved away from declarative to require the programmer to specify the low-level imperative steps required to perform batch tasks

2. KV store

- Only store KV pairs, no schema
- Generally fine if keys and values are flat and any processing required for the values can be done at app layer
- Downsides:
 - App layer will need to parse any complex V into individual fields, and also there are no secondary indexes.
 - No way to do joins or combine information across multiple KV pairs.
- But, there are still some popular (transactional) KV stores - thanks to their speed - DynamoDB is one
- Sometimes people write new DBs using existing KV store as storage engine
- From paper "*Many have either matured into RM systems or are only used for specific problems*"
- Most common examples: Redis, DynamoDb

3. Document DBs

- JSON data model - flexibility, performance (nesting helps keep info together, no join necessary)
- Different API that didn't involve SQL, and mostly used either a MR interface, or a bit closer to relational operators
- Very complicated to work with, optimization is hard for the system - onus is on user to write optimized "plans"
- Also same issue as hierarchical/nested data models - redundancies, limited physical data independence
- Likely will diminish over time - with JSON becoming a type in traditional RDBMSs
- Examples: MongoDB, Couchbase

4. Wide Column DBs

- A canonical implementation here was Google's BigTable in 2004, which got abandoned in 2010s, replaced by Spanner, fully relational
- Main idea was to have different subsets of rows with different subsets of columns
- Complicated to work with, similar issues as the union typing issues mentioned earlier

5. Text search engines

- Supporting text search by intersecting posting lists
 - essentially boolean conjunctive/disjunctive queries
- Specialized systems like Elastic, Solr, Lucene, ...
- But could also use a text index on a text column in a DBMS
 - May not be as performant but more convenient

6. Array databases

- Vectors, matrices, tensors,
 - Data where the "location" is paramount
- Main use-case is scientific, e.g., astronomy data
- Could also be used for LA and therefore ML but is an imperfect fit because usual expectation is that the arrays are sparse
- Some specialized systems, like SciDB
- But SQL now has array capabilities as of 1999

7. Vector databases

- A form of dense arrays, but storing embeddings from DL
- Main use-case is embedding similarity; k-NN
- Specialized systems Pinecone, Milvus, and Weaviate
- Most DBs have added vector capabilities as well
- Text search engines have also added them

8. Graph databases

- Most compelling use-case: primitives for graph traversal
- Performance may be better for RM with
 - Node (node_id, node_data)
 - Edge (node_id_1, node_id_2, edge_data)
- But ergonomics of PG or triple store languages are definitely better.
- Most prominent example: neo4j and Cypher

Big takeaway:

- RM + extensions (UDTs, UDFs) are enough to handle the vast majority of use-cases
- Most else is about special use cases (eg search) or applications (eg science)