

Lecture 4 - Berkeley History

Papers for today

- The POSTGRES Next-Generation Database System (1991) -
 - written close to when the Postgres team commercialized it.
 - full description of Postgres vision and architecture
- The Design of the Postgres Storage System (1987) -
 - a description of their storage manager using version and no-overwrite,
 - rather than WAL
- Looking Back at Postgres (1999) - retrospective on impact of Postgres
- Reason we're reading it is because Postgres did a "lot" of things - and had huge impact across a range of topics
 - Extensibility (types, functions, operators)
 - they call this "object management"
 - vast majority of this is done through existing system catalogs (an another example of dogfooding)
 - Rules system
 - they call this "knowledge management"
 - Multi-versioned storage
 - no overwrite storage - departure from full WAL
- Q: Thoughts on the paper?

CACM paper

Background

- Written in 1991; Postgres under development since 1986. Earliest papers:
 - The design of POSTGRES, SIGMOD'86
 - The POSTGRES Data Model, VLDB'87
- Wild that you can still run a piece of software 35 years later!

Design Philosophy

- *Orientation towards a set-oriented-query language*
 - POSTQUEL (PQ)
 - SQL support arrived with **Postgres95 (1994)**, and the system was renamed **PostgreSQL in 1996**
 - Some vestiges of CODASYL still remain (but deemphasized)
 - Declarative, set-oriented querying was central...
 - But: enable a "fast path"
 - Essentially allow functions to be registered and executed;
 - Can mix imperative code, PQ, as well as low-level functions ("get next record")
- *Multilingual access*
 - Departure from OODBs at the time that persist all program objects and were tied to a PL ->
 - here, you push PL constructs into a database
 - and instead use any PL to access it
 - this separation is now standard
- *a small number of concepts*
 - classes, inheritance, types, functions
 - of these inheritance is perhaps least emphasized in modern DBs, most of the others are folded into modern DBs in some way

- Honestly feels like a lot - but not compared to the complexity of systems at a time

Data Model

Classes and Inheritance

- a class \equiv relation; instances \equiv records
- Instances have OIDs - identifiers that can't be modified
- Inheritance:
 - `create EMP (name = c12, salary = float, age = int)`
 - `create SALESMAN (quota = float) inherits EMP`
- Can also inherit from multiple paths
- Different types of "classes"
 - Base/real - just a relation
 - Derived/virtual - just a view
 - Version - will see this later

Types

- Three types of types (ha)
 - Base types (which can also be extended)
 - Arrays of base types
 - Composite types
- Base types
 - The usual, but you can also extend to *Abstract Data Types*
 - Users can declare, they just need to provide a way to serialize it into strings
 - Though this may be costly/inefficient
 - `'create DEPT (dname = c10, floorspace = polygon, mailstop = point)'`
 - `replace DEPT (mailstop = "(10, 10)") where DEPT.dname = "shoe"`
 - Early form of extensible ADTs in modern DBs
- Arrays of base types
 - `create EMP (name = c10, salary = float[12])`
 - `retrieve (EMP.name) where EMP.salary[4] = 1000`
 - Most databases now support array operators
- Complex objects
 - Any class is also a type; and when you add that class as an attribute to another class, it becomes a container for zero or more instances of that class (weird, i know)
 - `create EMP (name = c10, salary = float[12], manager = EMP, coworkers = EMP)`
 - Both `manager` and `cworkers` end up being arbitrary sized lists of instances, essentially - afaict, no way of constraining to be a single instance
 - Q: how would we implement this in standard relational model?
 - Q: we learned about pitfalls of CODASYL - no nesting, no pointers - isn't this bad??
 - Yes, partially - even though nesting is present, it is not tied to physical implementation
 - Second flavor of complex object is a `set` - collection of arbitrary types
 - This feels just like JSON to me
 - `retrieve (EMP.manager.age) where EMP.name = "Joe"`
 - Dot notation, much like you would use in XPath, or in JSON in SQL

Functions

- Three types of functions: C functions, operators, PQ functions
- C UDFs:
 - arguments are base/composite types
 - `retrieve (DEPT.dname) where area (DEPT.floorspace) > 500`
 - `retrieve (EMP.name) where overpaid (EMP)`
 - Note: no optimization of such UDFs
 - still a difficult problem
- Operators:
 - `retrieve (DEPT.dname) where DEPT.floorspace AGT "(0,0), (1,1), (0,2)"`
 - Can be associated with access methods
 - R-Trees, KD-Trees, Joe later added GiST as part of his PhD thesis
 - Modern impact: Postgres uses GiST as a foundation (e.g., geometric, ranges); alongside other extensible index families (GIN, SP-GIST, BRIN).
- PQ functions:
 - Any collection of commands in PQ can be used to package a function
 - `define function high-pay returns EMP as retrieve (EMP.all) where EMP.salary > 50000`
 - Can use in PQ queries as well (essentially a subquery)
 - `append to EMP (name = "Sam", salary = 1000, manager = mgr-lookup ("shoe"))`
 - Or directly as well - this is that *fast path*
 - Generally user can directly call the parser, optimizer, executor, access methods, buffer manager, etc.
 - *it provides direct access to specific functions without checking the validity of parameters.* - therefore more efficient

Query Language Extensions

- Nested queries
 - Transitive closure
- These two are more interesting:
- Time travel
 - `retrieve (EMP.salary) from EMP [T] where EMP.name = "Sam"`
 - Inheritance
 - `retrieve (E.name) from E in EMP* where E.age > 40`
 - retrieves from all subclasses of `EMP` too

Rules System

- Goal - go beyond referential integrity and further push application logic and constraints to the database
 - *"referential integrity [7], which is merely a simple-minded collection of rules."*
- PG's rule system was very bold, encompassing:
 - view maintenance
 - triggers
 - integrity constraints
 - referential integrity
 - version management
- `on new EMP.salary where EMP.name = "Fred" then do replace E (salary = new.salary) from E in EMP where E.name = "Joe"`
- or equivalently `on retrieve to EMP.salary where EMP.name = "Joe" then do instead retrieve (EMP.salary) where EMP.name = "Fred"`
 - triggers on a read! Very unusual

- Former approach would work better if few updates of Fred, more reads of Joe latter if few reads of Joe, more updates of Fred.
 - not ideal that user has to pick
- Two main implementation methods:
 - Tuple level processing
 - Query rewrite
- Tuple level processing is very unusual
 - In example 1 above, marker placed on salary attribute of Fred's instance
 - If update touches attribute, then rule is called
 - Many corner cases; e.g.,
 - Fred's name is changed then Marker is dropped
 - Joe exists before Fred - marker added for first Fred
 - Performance issues with too many markers - escalate and add an enclosing marker
 - This only makes sense if there are specific rules for a small # of records
 - If rule touches most of table, then running this subroutine per query is overkill
- Query rewrite
 - Works better for queries that touch more
 - But rewriting while preserving semantics is nontrivial
 - Generally there is confusion around semantics with multiple instances - tuple level and query rewrite approach will do different things
- Modern triggers in practice go for more straightforward implementations: ECA - primarily for updates
- Rule systems for implementing views
 - `define view TOY-EMP (EMP.all) where EMP.dept = "toy"` implemented as:
 - `on retrieve to TOY-EMP then do instead retrieve (EMP.all) where EMP.dept = "toy"`
 - What is known at the time: updating to views is only done when it is unambiguous
 - Q: what is an example of a query where the view update is not ambiguous? ambiguous?
 - Postgres allows users to implement their own view update logic
- Rule systems for implementing versions:
 - `create version my-EMP from EMP`
 - Basically a fork, where updates can be done to this version without impacting the original
 - implemented as:
 - two classes:
 - `EMP-MINUS (deleted-OLD)`
 - `EMP-PLUS (all-fields-in EMP, replaced-OID)`
 - rewritten queries:
 - `on retrieve to my-EMP then do instead
 - `retrieve (EMP-PLUS.all)
 - `retrieve (EMP.all) where EMP.OID not in {EMP-PLUS.replaced-OID} and EMP.OID not in {EMP-MINUS.deleted-OID}`

Storage Engine

- Also the focus of the second paper!
- wanted "*to do something different*" - thought WAL was overkill
- WAL enables **STEAL + NO-FORCE** buffer management:
 - pages can be written before commit (steal)
 - pages don't have to be forced at commit (no-force)
 - instead, force the log (sequential I/O) to guarantee durability

- ARIES wasn't yet known (1992) - but mostly it focused on "completing" the story when there are crashes during recovery etc.
- Fundamentally different approach for recovery
 - a "no overwrite" storage manager
 - so: purpose served by WAL is done in regular storage
 - WAL now replaced by a log of status of each transaction (2 bits)
 - tuple versions + transaction status log is sufficient!
- Benefit 1: instantaneous abort and crash recovery
 - No need to undo updates of uncommitted txns, redo updates of committed txns
 - Everything is reflected in the regular data - we just need to abort the inprogress txns, and not change the regular data
- Benefit 2: time travel is possible.
 - Queries for data "as of" time T
 - Theoretically one could try to do so with the log but the log would need to be augmented.
- Downside:
 - All pages for a committed transaction must be forced to disk
 - to ensure durability of committed transaction
 - this can involve many random writes;
 - contrast to flushing a log, which is sequential & cheaper
 - generally this would work well with a more "stable" memory
 - Modern context:
 - Lots of hope for NVRAM - but never materialized
 - most prominent recent commercial attempt—**Intel Optane persistent memory**—was discontinued in 2022
 - "All data" mixed in; if there are lots of aborts or even lots of updates, there's lots of data to wade through
- Postgres can also migrate from one long term storage format (disk) to archival (tape)

Additional Details from second paper

- Lots of "back of envelope" math to justify the fact that transaction statuses could be stored in main memory
 - But now not a big deal - memory is much cheaper and plentiful!
- Per record, store the following:
 - OID a system-assigned unique record identifier
 - Xmin the transaction identifier of the interaction inserting the record
 - Tmin the commit time of Xmin (the time at which the record became valid)
 - Cmin the command identifier of the interaction inserting the record
 - (similarly, Xmax, Tmax, Cmax)
 - PTR a forward pointer
- Basically Xmax, Tmax, Cmax will be NULL while the record version is still valid, and then get updated when the record is updated or deleted.
- For new versions of a record, can additionally only restrict to storing delta (only attributes that change)
- A record is valid at time T if the following is true:
 - Tmin < T and Xmin is a committed transaction and either:
 - Xmax is not a committed transaction or
 - Xmax is null or
 - Tmax > T
- Uses locking

Some perspectives from Joe's paper

- Abstract Data Types and Functions:
 - All major vendors have support for hierarchical data via JSON/XML, using very similar path notation to what Postquel did
 - All major vendors support UDFs
 - One could argue that MR revolution was a need to support UDFs + parallelism -> now certainly looks that way
 - Extensible data types
 - Postgres now supports GIN (search for unstructured data), GiST
 - OR won out against OO - additional OO functionality came in the form of ORMs - which were different from both
 - UDF query optimization
 - Still really hard
 - Q: Suppose we have a user defined function predicate - should we still push it down?
 - Joe worked on this problem as part of his thesis
- Active databases
 - Triggers are part of the database standard
 - But are really hard to enforce in practice
 - See similar ideas in materialized view maintenance, CEP, streaming
- Version-centric storage
 - Realization: current state is simply a view over versions
 - Modern incarnations: LSM trees in KV stores
 - Unfortunately never excelled in performance, versioning and time-travel was replaced by WAL in 2001; see announcement (time travel was removed even earlier, 1998)
 - But - within postgres, multi-versioning ended up being useful to support MVCC (added in 1999) - and snapshot isolation - more on that at some future point
- Open Source
 - SQL porting: 1995 - called PostgreSQL after that
 - Many startups, often around parallel postgres
 - Andy Pavlo's 2026 seminar series [Postgres vs. the world](#)

"Postgres was designed for extensibility, and that design was sound. With extensibility as an architectural core, it is possible to be creative and stop worrying so much about discipline: you can try many extensions and let the strong succeed."

Another lesson is that a broad focus—"one size fits many"—can be a winning approach for both research and practice.