

Lecture 8 - Isolation Levels Revisited

- Two papers:
 - Adya's Generalized Isolation Level Definitions
 - A critique of SQL Isolation Levels
- Goal: begin exploration of weaker isolation levels beyond the original degrees of consistency paper

To discuss in class:

- Q: How do folks feel about the project?
- Q: How are the readings treating you?
- Announcement: special seminar
- Q: How do we make up the lecture?

Prehistory and a Spooky Problem

- Original Gray proposal:
 - degree 0 - short duration write locks
 - degree 1 - long duration write locks
 - degree 2 - long duration write locks, short duration read locks
 - degree 3 = serializable - long duration write locks, long duration read locks
- Unfortunately, degree 3 doesn't actually guarantee serializability, except under a simplified model of locking of individual items
 - as opposed to for SQL
- Phantoms come from predicates/ranges, not individual items:
 - Suppose T1 reads all rows where $\text{Salary} > 100K$
 - Then T2 inserts a new row with $\text{Salary} = 120K$ and commits
 - If T1 repeats the predicate, it sees an extra row: a *phantom*
- The paper that formalizes this came right after the original degrees of consistency paper: [The Notions of Consistency and Predicate Locking](#)
 - Basically, degree 3 is insufficient for serializability
 - Instead, proves that predicate locking + 2PL \Rightarrow serializability
 - What is predicate locking: basically "lock" corresponding to a predicate, e.g., $\text{Salary} > 100K$
 - T2's insertion here would not be allowed as it conflicts with the predicate
- Q: Easy enough, but how to enforce? Why is locking predicates hard?
 - Really difficult in practice:
 - General lock manager for predicates is not practical
 - How does know if $\text{Salary} > 100K$ overlaps with $\text{Dept} = \text{EECS}$?
- Modern consensus on how to handle phantoms:
 - Restrict predicate locking to indexable ranges (e.g., B+tree intervals)
 - Lock ranges e.g., as in [Lomet 1993](#)
 - Operationally, easy to coordinate in a hierarchical index (Q: what does this remind you of? remember multi-granularity locking?) - details differ
 - Or: if no index exists, can use coarser grained locking (page/table)
 - Or: not a locking-based approach at all (e.g., SI - more next or SSI - will see next time)

ANSI Isolation Levels and Problems

- This plus refinement by Date, led to ANSI isolation levels SQL 92


- Goal - implementation independence
 - Actual impact: Definitions less constraining
- Described *phenomena* that are proscribed, such as dirty read, non-repeatable read, phantoms, ...
 - And corresponding isolation levels, read uncommitted, read committed, repeatable read, serializable, etc.
 - We'll see what these are in a bit
- Berenson and co said:
 - These are ambiguous, and do not exclude some anomalous behavior
 - No clear mapping between locking isolation levels and the phenomena
 - Also doesn't identify other popular isolation levels
- So redefined the phenomena:
 - P0: w1[x] ... w2[x] ... (c1 or a1) - dirty write
 - Can't write an object if an active transaction has already written to it.
 - missing in the ANSI-SQL definition
 - P1: w1[x] ... r2[x] ... (c1 or a1) - dirty read
 - Can't read an object an active transaction has already written
 - P2: r1[x] ... w2[x] ... (c1 or a1) - non-repeatable read
 - Can't write an object that has been read by an active transaction
 - P3: r1[P] ... w2[y in P] ... (c1 or a1) - phantoms
 - Can't modify rows matching a predicate that an active transaction has read

Locking Isolation Level	Proscribed Phenomena	Read Locks on Data Items and Phantoms (same unless noted)	Write Locks on Data Items and Phantoms (always the same)
Degree 0	none	none	Short write locks
Degree 1 = Locking READ UNCOMMITTED	P0	none	Long write locks
Degree 2 = Locking READ COMMITTED	P0, P1	Short read locks (on both)	Long write locks
Locking REPEATABLE READ	P0, P1, P2	Long data-item read locks, Short phantom read locks	Long write locks
Degree 3 = Locking SERIALIZABLE	P0, P1, P2, P3	Long read locks (on both)	Long write locks

- Other phenomena: P4: r1[x]... w2[x] ... w1[x]... c1 - lost update
 - Note that P4 is already proscribed by P2 - a special case of P2
 - So helps identify isolation levels between READ COMMITTED and REPEATABLE READ
- Two new serializable levels, all between Degree 2 and Degree 3
 - Cursor Stability
 - Details are not important but prevents P4
 - Amounts to READ COMMITTED << Cursor Stability << REPEATABLE READ
 - Snapshot Isolation
 - MVCC approach not captured in ANSI Levels
 - Txns take and operate on copies, like OCC
 - Txn Ti are allowed to commit if no other txn in the interim wrote data that Ti wrote
 - "First committer wins" on ww conflicts
 - Like cursor stability, lost update (P4) can't happen because w2 above won't be able to commit
 - But not serializable:
 - r1[x=50] r1[y=50] r2[x=50] r2[y=50] w1[y=-40] w2[x=-40] c1 c2
 - Essentially: r1[x]...r2[y]...w1[y]...w2[x]...(c1 and c2 occur)
 - this is called *write skew* - there is an analogous one called *read skew*

- We can show that READ COMMITTED ≪ Snapshot Isolation
- But, interestingly REPEATABLE READ ≪> Snapshot Isolation
 - For example, write skew is prohibited in REPEATABLE READ
 - Likewise, reading the same item cannot result in different values in SI
- Read consistency
 - Omitted details (a bit hard to follow from Berenson)
 - READ COMMITTED ≪ Read Consistency ≪ REPEATABLE READ
 - and Read Consistency ≪> Cursor Stability

Serializable = Degree 3 = {Date, DB2} Repeatable Read



Onto Adya

- But P0 is essentially a long-duration write lock!
 - P1 requires long-duration write locks, and readers to take short-duration read locks
 - P2 implies long-duration read and write locks
 - P3 involves predicate locks
- Problem: P0, P1, P2, P3 disallows many legal schedules
 - P0/P2 - can happen in optimistic implementations - many txns are editing local or reading copies simultaneously, but some will be forced to abort
- At its core:
 - Preventative approach expresses phenomena in terms of single-object histories
 - both OCC and MVCC-based approaches can't be modeled that way
 - Preventative approach also has same properties for running and committed txns
 - but OCC-based approaches have weak guarantees while running, and then strong consistency for committed txns
- What Adya hopes for: Implementation-independent spec of isolation levels
 - Both correct (rule out bad histories) and complete (include all conflict serializable schedules)


- For lower isolation levels, more permissive than Berenson
- Done by proscribing (eliminating) different types of cycles in a serialization graph

Adya's Formalization

- $x_{i,1}, x_{i,2}, \dots$ denotes k th modification of x by T_i ; final modification is x_i .
- After committing T_i installs x_i .
- Each object starts in the x_{init} state (unborn), then becomes "visible", then becomes "dead"
 - If comes back again, essentially a new object
- $w_i(x_{i,m})$ - write of x by T_i .
- $r_j(x_{i,m})$ - read by T_j of the m th version written by T_i .
- A history is a partial order that:
 - preserves individual op order by transactions
 - for something to be read, it must have been written - $r_j(x_{i,m})$ must be preceded by a $w_i(x_{i,m})$
 - if a txn writes then reads immediately without an intervening event, then it must see its write, e.g., $r_i(x_j)$ after a $w_i(x_{i,m})$ should be $j = i, m$
 - must be complete: if you see any events for T_i it must be committed
 - note that we don't reason about uncommitted txns
 - i.e., there can be violations that "get fixed later"
 - one way to reason about partial histories is to add aborts for uncommitted txns
 - "worst-case behavior"
- In addition to events, we want a total order \ll of versions of object created by committed txns (also called committed versions)
 - $x_{init} \dots$ (various visible versions) $\dots x_{dead}$
 - if $r_j(x_i)$ is in a schedule, it is visible
- Schedule:
 - $w_1(x_1)w_2(x_2)w_2(y_2)c_1c_2r_3(x_1)w_3(x_3)w_4(y_4)a_4 \dots x_2 \ll x_1$
 - x_2 serialized before x_1 even though commits are other way around
- Predicates
 - $VSet(P)$ set of versions of all tuples read by P
 - $r_i(P : VSet(P))r_i(x_j)r_i(y_k) \dots$ - where x_j, y_k all versions in $VSet$ that match P and read by T_i
 - While most of the real-estate in 4.3 onwards is on predicates, we are going to ignore it and try to focus on the simpler setting without predicates

Direct Serialization Graphs

- Three types of edges in the graph:
 - direct write dependency $ww : T_i$ installs x_i and T_j installs next version
 - direct read dependency $wr : T_i$ installs x_i and T_j reads x_i or (predicate based version)
 - direct anti dependency $rw : T_i$ reads x_h and T_j installs next version (or predicate based version)
- History : $w_1(z_1)w_1(x_1)w_1(y_1)w_3(x_3)c_1r_2(x_1)w_2(y_2)c_2r_3(y_2)w_3(z_3)c_3$ plus
 - $x_1 \ll x_3, y_1 \ll y_2, z_1 \ll z_3$



- Serializable as $T_1 \ll T_2 \ll T_3$
- Definitions:
 - *direct depends*: if there is a direct read or write
 - *depends*: if there is a path

Isolation level PL-1

- Prohibiting something analogous to
 - P0: $w1[x] \dots w2[x] \dots (c1 \text{ or } a1)$ -- dirty write
- The generalizable phenomenon, G0 is *no write dependency cycles*
- $w_1(x_1, 2)w_2(x_2, 5)w_2(y_2, 5)c_2w_1(y_1, 8)c_1$ alongside
 - $x_1 \ll x_2, y_2 \ll y_1$
- $T_1 \rightarrowtail \text{ww} \rightarrowtail T_2$ (and back ww edge)
- More permissive because same object can be concurrently modified unlike P0


Isolation level PL-2

- Prohibiting something analogous to
 - P1: $w1[x] \dots r2[x] \dots (c1 \text{ or } a1)$ -- dirty read
- This is a bit complicated (omitting predicate versions):
 - G1a - No *aborted reads*
 - $w_1(x_{1,i}) \dots r_2(x_{1,i}) \dots (a_1 \text{ and } c_2 \text{ in any order})$
 - Can't have a committed txn reading a version created by an aborted one
 - G1b - No *intermediate reads*
 - $w_1(x_{1,i}) \dots r_2(x_{1,i}) \dots w_1(x_{1,j}) \dots c_2$
 - Can't have a committed txn reading "intermediate" versions
 - Together G1a and G1b ensures committed txns only read objects that existed or will exist at some instant in the committed state
 - G1c - No *circular info flow*
 - Directed cycle of dependency edges (read or write edges)
 - Also includes G0 by definition
- G1 is weaker than P1 because G1 allows reads from uncommitted txns
- And lock-based impl (long write locks, short read locks) will guarantee G1
 - G1a and G1b can't happen because T_2 can't read before T_1 commits
 - G1c if T_1 read something T_2 wrote, then T_2 can't read/write something T_1 wrote

Isolation Level PL-3

- Prohibiting something analogous to:
 - P2: $r1[x] \dots w2[x] \dots (c1 \text{ or } a1)$ -- non repeatable read
- Simple enough: G2: No cycles with one or more anti-dependency edges.

- G2 + G1 together give something analogous to P2
- If we omit predicate centric dependency edges, we have the variant PL-2.99
- G2 also prevents lost update (just like P2 prevents P4): from Adya's Thesis
 - Lost update history : $r1(x_0, 10) r2(x_0, 10) w2(x_2, 15) c2 w1(x_1, 14) c1$
 - with $x_0 \ll x_2 \ll x_1$



More reading

- Adya's thesis also formalizes SI and Cursor Stability through the DSG, and identifies new isolation levels as well